

# Deep Neural Networks

## Vision

hands on session

Alexander Boettcher

# Topics

- Introduction
- DNN
- CNN

# Introduction - Deep Learning preconditions

- Theory

Already developed starting in the 1970s

- Data

Available since digital cameras and the internet spread

- Processors

Parallelization in GPUs is available

- Frameworks

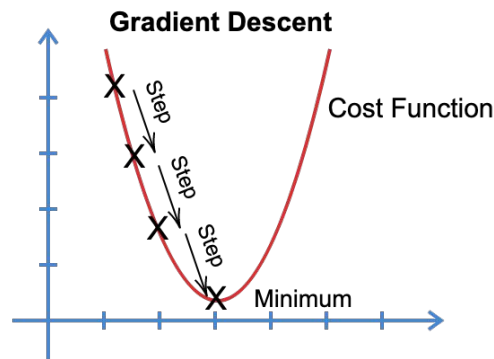
Popular frameworks currently Pytorch and Tensorflow

# Introduction - Theory

- Mathematical numerical optimization


$$\underset{x}{\text{minimize}} \quad f(x)$$

- First order optimization (e.g. gradient descent)
  - In DL it is easy to obtain derivative (simple basic functions and chain rule)
- Fast optimization is key
  - Aim: reach the (local) optimum using few evaluations of the function and its derivative



# Introduction - Data

- Sets of images
- Stored in matrixes (or tensors)
- Common shape of a data tensor (with four axes):
  - number of images,
  - number of channels,
  - number of columns,
  - number of row



# Introduction - Processors

- GPUs (Graphics processing units) are popular in Deep Learning
  - Specialized on linear algebra
    - matrix multiplication
    - Convolutions
  - Many cores for parallelization of matrix operations.
  - Most operations during optimization are matrix operations!

# Introduction - Frameworks

- Important features of deep learning frameworks
  - Automatic differentiation
  - Set of popular optimizers
  - Popular (pretrained) networks

# Introduction - Summary

## Takeaways:



- Deep learning is based on **mathematical numerical optimization**
- The workhorse of deep learning are **optimizers** (e.g. gradient descent)
- Deep learning frameworks calculate the gradient using **automatic differentiation**
- Deep learning **libraries** provide everything needed to do all of the above at ease.



# Introduction - Hands on

Jupyter Notebook: Session\_1\_introduction


# DNN - Neuron




```
import torch
class Neuron:
    def __init__(self):
        self.w_0 = torch.autograd.Variable(torch.Tensor([1]), requires_grad=True)
        self.w_1 = torch.autograd.Variable(torch.Tensor([1]), requires_grad=True)
        self.bias = torch.autograd.Variable(torch.Tensor([1]), requires_grad=True)

    def f(self, x_0, x_1):
        linear = x_0*self.w_0 + x_1*self.w_1 + bias
        non_linear = np.clip(linear, 0., np.inf)
        return non_linear


n = Neuron()
```



(a)





(b)



(c)

# DNN - Neural Network



```
import torch
class Neuron:
    def __init__(self):
        self.w_0 = torch.autograd.Variable(torch.Tensor([1]), requires_grad=True)
        self.w_1 = torch.autograd.Variable(torch.Tensor([1]), requires_grad=True)
        self.bias = torch.autograd.Variable(torch.Tensor([1]), requires_grad=True)


    def f(self, x_0, x_1):
        linear = x_0*self.w_0 + x_1*self.w_1 + bias
        non_linear = np.clip(linear, 0., np.inf)
        return non_linear

n_0 = Neuron()
n_1 = Neuron()
n_2 = Neuron()


def network(x_0, x_1, x_2, x_3):
    y = n_0(n_1(x_0, x_1), n_2(x_2, x_3))
    return y
```

# DNN - Neural Network Architectures

Feed-forward networks  
(no loops allowed)



Recurrent networks  
(loops allowed)



# Introduction - Hands on

Jupyter Notebook: Session\_2\_neural\_networks

# Tasks

- Classification

Out of a given set of classes  $y$  predict the most likely class

- Regression

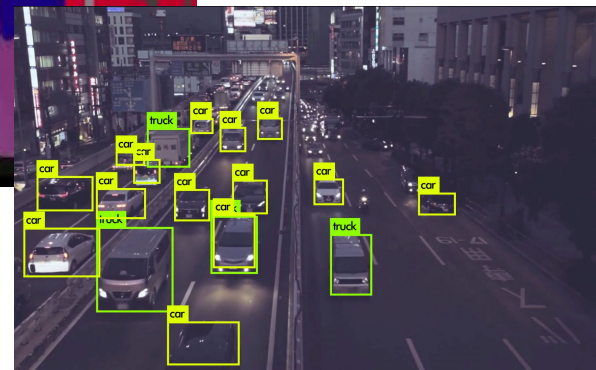

$y$  predicts a floating point number

- Segmentation


Out of a given set of classes  $y$  predicts the most likely class for each pixel in the input image

- Detection

Out of a given set of classes  $y$  predicts boxes for each instance of the classes in the input image



# Classification - Example




Example: classify if the input image shows a car

- no car:  $y_0 = 0$
- car:  $y_0 = 1$

How to make  $y_0$  to return binary classification results?

# Classification - Example




Example: classify if the input image shows a car

- no car:  $y_0 = 0$
- car:  $y_0 = 1$

Use a sigmoidal non-linearity in the last layer



# Classification - Example




Example: classify if the input image shows a car

- no car:  $y_0 = 0$
- car:  $y_0 = 1$

Random output due to random initialization of the network

# Classification - Example




Example: classify if the input image shows a car

- no car:  $y_0 = 0$
- car:  $y_0 = 1$


Optimize the difference  
(loss) between the output  
and the label

# Classification - Cross entropy loss

$\text{loss}(\text{model}(x), \text{label}=1)$



# Classification - cross entropy loss



# Introduction - Hands on

Jupyter Notebook: Session\_3\_neural\_networks\_optimization


# Classification - Example



Iris Versicolor

Iris Setosa

Iris Virginica



Example network with

- 4 inputs
- 3 output classes

Classes (one hot encoding):

- Iris Setosa: [1, 0, 0]
- Iris Versicolour: [0, 1, 0]
- Iris Virginica: [0, 0, 1]

# Pytorch

As the examples get more complicated, let's use Pytorch to specify the model and load the data.

# Introduction - Hands on

Jupyter Notebook: Session\_4\_DNNs



# CNNs

Natural images are spatially invariant:


- The statistics is over all the space the same
- A neuron that detects an edge conducts the same task no matter at what spatial location

The correlation of pixel intensities decrease the further away the pixels are:

- Closeby pixels are related e.g. an edge might affect several collocated pixels
- Pixels far apart have almost no relation to each other

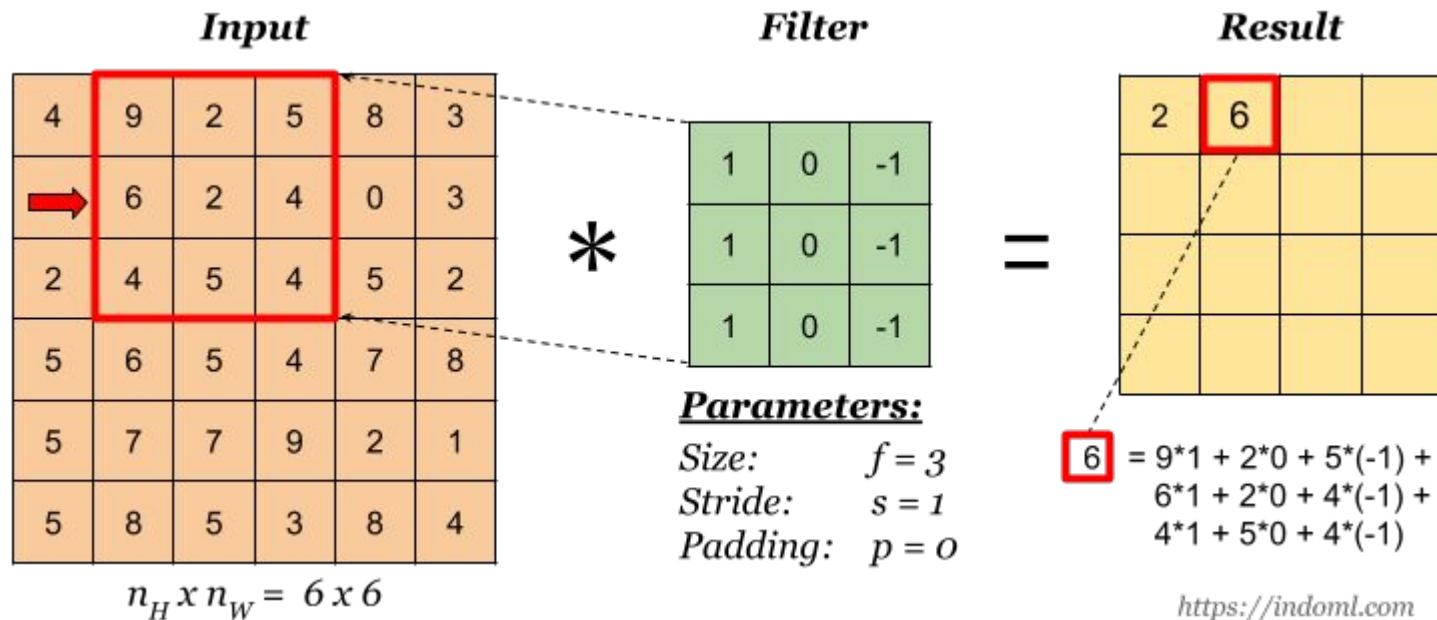
# CNNs

Fully connected layers connect all neurons of the lower layer with all layers of the upper layer:



# CNNs

Fully connected layers connect all neurons of the lower layer with all layers of the upper layer:



# Introduction - Hands on



Jupyter Notebook: Session\_5\_CNNs

# Finetuning



Learning big networks from scratch

- Demands large amount of data
- Takes long time
- Demands fast GPUs

It is easier to reuse a trained network and only retrain a part of it.



# Finetuning



# Introduction - Hands on

Jupyter Notebook: Session\_4\_finetuning